

```
import tensorflow as tf
import kagglehub
import shutil
import os
from google.colab import drive
from tensorflow.keras.models import Sequential, load_model

#get latest version of dataset
path = kagglehub.dataset_download("seanscully29/flowers-classification")
print("Path to dataset files:",path)
```

➦ Path to dataset files: /kaggle/input/flowers-classification

```
local_path="./flower_species_dataset"
os.makedirs(local_path, exist_ok=True)
```

```
try:
    for item in os.listdir(path):
        s = os.path.join(path,item)
        d = os.path.join(local_path,item)
        if os.path.isdir(s):
            shutil.copytree(s,d)
        else:
            shutil.copy2(s,d)
except FileExistsError: #If code accidentally runs again, print text signaling that the path already exists instead of throwing an error
    print(f"{local_path} already exists.")
```

```
print(f"Dataset saved at: {local_path}")
```

➦ Dataset saved at: ./flower_species_dataset

```
source_path="/content/flower_species_dataset/flowerdataset"
training_path='./flower_species_dataset/training'
testing_path='./flower_species_dataset/testing'
classes=["black_eyed_susan","calendula","california_poppy","coreopsis","iris"]

model_path="flower_dataset_model.h5"
```

```

for cls in classes: #creates subdirectories for each flower class
    os.makedirs(os.path.join(training_path,cls),exist_ok=True)
    os.makedirs(os.path.join(testing_path,cls),exist_ok=True)

for cls in classes:
    files=os.listdir(os.path.join(source_path,cls)) #Lists image files in source directory

    #split 65% of the images to the training set, and 35% to the testing set
    split=int(0.65 * len(files))
    training_files=files[:split]
    testing_files=files[split:]

    for t in training_files: #Adds image files to recently created subdirectories on a 65/35 split
        shutil.move(os.path.join(source_path, cls, t),os.path.join(training_path,cls,t))
    for t in testing_files:
        shutil.move(os.path.join(source_path, cls, t),os.path.join(testing_path,cls,t))

# Define ImageDataGenerator for augmentation and rescaling
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_directory='/content/flower_species_dataset/training'
test_directory='/content/flower_species_dataset/testing'
train_datagen = ImageDataGenerator(
    rescale=1./255, #Normalizes pixel values
    rotation_range=20, #Randomly rotates image by up to +- 20 degrees
    width_shift_range=0.2,
    height_shift_range=0.2, #Shifts height/width by up to 20%
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

test_datagen = ImageDataGenerator(rescale=1./255)

# Load training dataset
train_generator = train_datagen.flow_from_directory(
    train_directory,

```

```
target_size=(150, 150),
batch_size=32,
class_mode='categorical'
)
```

Load testing dataset

```
test_generator = test_datagen.flow_from_directory(
    test_directory,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical'
)
```



```
Found 3314 images belonging to 5 classes.
Found 1787 images belonging to 5 classes.
```

```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
```

```
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(2,2),

    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(2,2),

    Conv2D(128, (3,3), activation='relu'),
    MaxPooling2D(2,2),

    Conv2D(256, (3,3), activation='relu'),
    MaxPooling2D(2,2),

    Flatten(),
    Dense(512, activation='relu'),
    Dense(5, activation='sigmoid') # Categorical with 5 classes
])
```

Compile the model

```
model.compile(loss='categorical_crossentropy',
```

```
optimizer='adam',  
metrics=['accuracy'])
```

```
➞ /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input`  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
history=model.fit(train_generator,epochs=15,validation_data=test_generator)  
model.save(model_path)
```

```
➞ /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class s  
self._warn_if_super_not_called()
```

Epoch 1/15

104/104 ————— 33s 252ms/step - accuracy: 0.4245 - loss: 1.2578 - val_accuracy: 0.4835 - val_loss: 1.2053

Epoch 2/15

104/104 ————— 22s 213ms/step - accuracy: 0.5368 - loss: 0.9910 - val_accuracy: 0.6435 - val_loss: 0.8580

Epoch 3/15

104/104 ————— 23s 221ms/step - accuracy: 0.6340 - loss: 0.8368 - val_accuracy: 0.7057 - val_loss: 0.7237

Epoch 4/15

104/104 ————— 22s 212ms/step - accuracy: 0.6934 - loss: 0.7286 - val_accuracy: 0.7146 - val_loss: 0.6978

Epoch 5/15

104/104 ————— 21s 205ms/step - accuracy: 0.7103 - loss: 0.7077 - val_accuracy: 0.7096 - val_loss: 0.6626

Epoch 6/15

104/104 ————— 22s 214ms/step - accuracy: 0.7321 - loss: 0.6388 - val_accuracy: 0.7555 - val_loss: 0.6183

Epoch 7/15

104/104 ————— 22s 214ms/step - accuracy: 0.7490 - loss: 0.6056 - val_accuracy: 0.7443 - val_loss: 0.6030

Epoch 8/15

104/104 ————— 22s 215ms/step - accuracy: 0.7512 - loss: 0.5945 - val_accuracy: 0.7499 - val_loss: 0.6195

Epoch 9/15

104/104 ————— 22s 207ms/step - accuracy: 0.7743 - loss: 0.5307 - val_accuracy: 0.7874 - val_loss: 0.5308

Epoch 10/15

104/104 ————— 21s 206ms/step - accuracy: 0.7811 - loss: 0.5321 - val_accuracy: 0.7773 - val_loss: 0.5597

Epoch 11/15

104/104 ————— 22s 214ms/step - accuracy: 0.7917 - loss: 0.5050 - val_accuracy: 0.7734 - val_loss: 0.5384

Epoch 12/15

104/104 ————— 23s 220ms/step - accuracy: 0.7907 - loss: 0.4988 - val_accuracy: 0.7957 - val_loss: 0.4975

Epoch 13/15

104/104 ————— 22s 210ms/step - accuracy: 0.8047 - loss: 0.4650 - val_accuracy: 0.8008 - val_loss: 0.4701

Epoch 14/15

104/104 ————— 22s 209ms/step - accuracy: 0.8288 - loss: 0.4249 - val_accuracy: 0.7678 - val_loss: 0.6114

Epoch 15/15

104/104 ————— 21s 204ms/step - accuracy: 0.7987 - loss: 0.4826 - val_accuracy: 0.8041 - val_loss: 0.4661

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is consid

```
import matplotlib.pyplot as plt
```

```
# === Visualization of training ===
```

```
print("Visualizing training results...")
```

```
# Plot accuracy
```

```
plt.figure(figsize=(12,5))
```

```
plt.subplot(1,2,1)
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
```

```
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.title('Model Accuracy')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Accuracy')
```

```
plt.legend() #Adds necessary labeling to plot
```

```
# Plot loss
```

```
plt.subplot(1,2,2)
```

```
plt.plot(history.history['loss'], label='Training Loss')
```

```
plt.plot(history.history['val_loss'], label='Validation Loss')
```

```
plt.title('Model Loss')
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Loss')
```

```
plt.legend()
```



```
plt.tight_layout()
```

```
plt.show()
```





Visualizing training results...



```
from tensorflow.keras.models import load_model
from PIL import Image, ImageOps
import numpy as np
import matplotlib.pyplot as plt
```

```
def predict_new_image(img_path):
    model = load_model(model_path)
```

```
    # Load image in color (RGB)
    img = Image.open(img_path).convert("RGB")
    img = img.resize((150, 150))
```

```
img_array = np.array(img).astype("float32") / 255.0
img_array = img_array.reshape(1, 150, 150, 3)
```

```
prediction = model.predict(img_array)
predicted_index = np.argmax(prediction)
confidence = prediction[0][predicted_index]
```

```
# Display
```

```
plt.imshow(img)
```

```
plt.axis("off")
```

```
plt.title(f"Prediction: {classes[predicted_index]}, Confidence: ({confidence:.2%})")#Prints name of classes most likely to be the flower :
```

```
plt.show()
```

```
predict_new_image("/content/calendula.jpg")
```



WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train the model.
WARNING:tensorflow:6 out of the last 6 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7f8b1c1c1c1c> were successful (the first call was not).

Prediction: calendula, Confidence: (94.02%)

